

Replica" and sharding

- 1) A dB is critical to sys.
- 2) ∵ sys's performance \propto dB's performance
- 3) $\begin{array}{c} \text{dB} \\ \text{eg 1) unavail.} \end{array} \quad \begin{array}{c} \text{sys} \\ \text{unavail.} \end{array}$
 $\begin{array}{c} \text{eg 2) low throughput} \\ \text{low throughput} \end{array}$
- 4) While des. in a sys and making it performin, we gotta ensure its dB is performant, else sys will fall apart

I]

Replica"

The act of duplicating data from 1 dB servers to others

use case (i) : To act as backup if main dB server fails

- ① ~~When~~ When main dB goes down, we can't perform any opns. (read, write, etc.) on the data ^{clients}
- ② To prev. clients from experiencing this we establish a secondary dB (backup) i.e. a replica of the main dB.

- ③ For this to happen the main dB, on receivin req.

→ performs whatever opn (read, write, etc) the req. is for

→ updates the replica s.t. replica == main dB
 i.e. in 'sync' with dB at any given pt. in time

④ we do this to ensure that replica dB can take over in case main dB fails

⑤ main dB fails → replica takes over as your main dB after some time ↓

now, they swap roles, if main dB takes over again as main dB

now, main dB is updated by replica

main dB comes back up

replica goes back to act as a standby.

⑥ write oprns to main dB take a lit longer as you gotta write data in main dB and replica to ensure replica is always 'in sync' (i.e. up-to-date) w.r.t. main dB.

Each write oprn of main dB should sync'dly be done in replica. If write oprn fails on replica due to some reason (e.g.: network part), then the write oprn. should also not be completed on main dB.

⑦ Replica should NEVER have stale data.

⑧ Redundancy of sys ↑

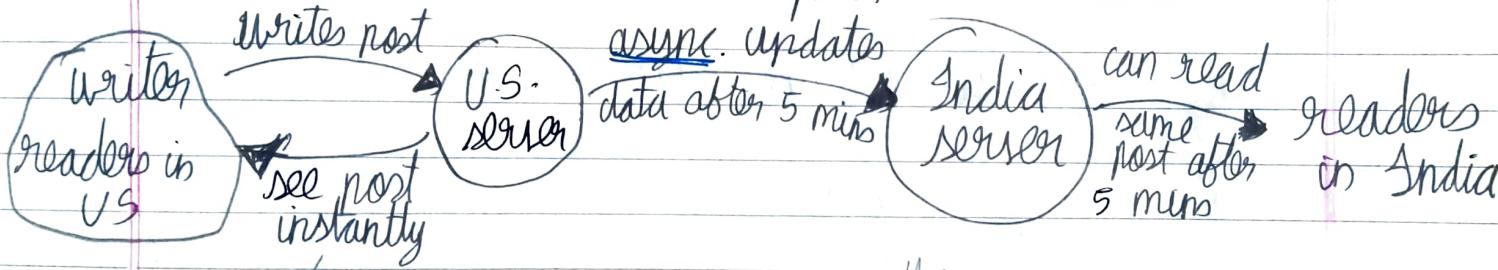
use Case (ii) : To move data closer to clients, thereby ↓ latency of accessin ~~and~~ specific data.

① Consistency : When you save data to a dB, it gets

stored but time taken \propto dist. of server from user.

- ② To bring server closer to user i.e. ↓ latency, we ~~can~~ use a replica server close to user and user stores his/her writes to it. (takes lesser time as server is closer)
- ③ We async.ly update data (say once/5min) present in all servers to ~~be~~ have consistent data everywhere
- ④ Used only in sys.s where we can afford the delay this async. update causes

e.g.: - 1 user in US writes post on LinkedIn. 1 user in India reads this post



Reason \rightarrow server is closer (not like cases where server is on other side of world so you gotta make a time consuming round trip to 1st write data (writer) and then read data (reader))

NOTE: vice versa Indian writers' posts ~~can't~~ are also read by US readers 5 mins after post is published

e.g.: - If you don't need result of a deployment of new feature in your app to be seen all over the world instantly.

II] Sharding

① It may happen that main dB is handling way more say (say, 1 billion/sec) than it should be and thus, a ~~bottleneck~~ bottleneck starts forming at main dB servers.

horizontally scaling of a server

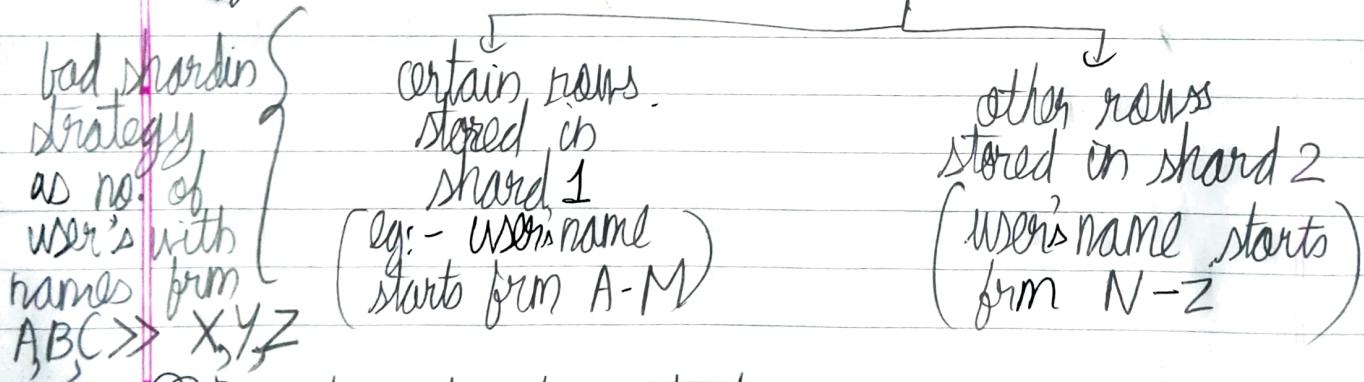
② So he ~~will split up main dB data~~ small nos. of data ~~is main~~ and store them in separate small servers (known as **shards**) and use them to ~~provide~~ provide data to diff. kinds of reqs.

③ Shards a.k.a. "data partitions"

Sharding : Act of splitting a dB into 2 or more pieces (**shards**)

Advantage : Throughput ↑ (as no bottleneck formed)

e.g.: - say we got an rdB → table



④ Popular sharding strategies

i) Based on clients' regions

ii) Based on type of data being stored

e.g.: - data → split → shard 1 → user data

iii) Sharding based on the hash of a column (only done for structured data)

- ⑤ Hotspots :- Certain shards that get more traffic than other shards cuz of the kinda data they store
- Cause → shardin key or hashin fⁿ is sub-optimal
- Effect → worsen distri. of workload across shards

⑥ Minimizing hotspots i.e. evenly splitton up data

i) Use a hashin fⁿ (H.F.) that guarantees uniformity in determinin what data goes to which shard.

ii) Also, we need to ensure ~~the~~ H.F. doesn't change much to ensure that a particular type of data always goes to the same servers

iii) Consistent hashin is kinda useful (H.F. is const)

→ Addin a new shard → ~~the~~ consistent hashin minimizes the pieces of data we gotta migrate from existin shards to new shard

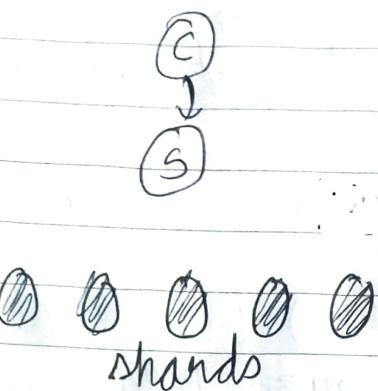
→ Shard goes down → consistent hashin can't do anything

↳ soln:- replicate each shard.

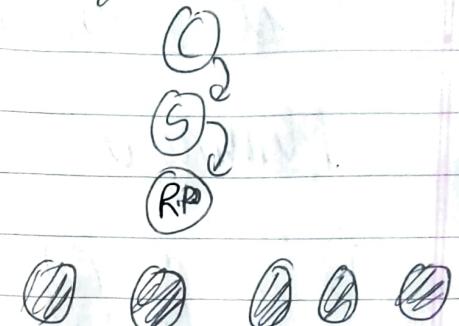
∴ ~~think~~ Give a lotta thought to comin up with a good logic to split up data while designin a sys.

(7) Storage of sharding strategy logic

M-1 → stored in server



M-2 → stored in reverse proxy b/w server & shards



Place of storage basically decides that based on clients req which shard is it → gonna req. data from
 OR
 write data to

```
//REVERSE PROXY - STORES LOGIC TO SPLIT INCOMING DATA (WRITE REQUEST) INTO SHARDS
```

```
const axios = require('axios');

const express = require('express');

const SHARD_ADDRESSES = ['http://localhost:3000', 'http://localhost:3001'];

const SHARD_COUNT = SHARD_ADDRESSES.length;

function getShardEndpoint(key) {

  const shardNumber = key.charCodeAt(0) % SHARD_COUNT;

  const shardAddress = SHARD_ADDRESSES[shardNumber];

  return `${shardAddress}/${key}`;

}

app.post('/:key', (req, res) => {

  const shardEndpoint = getShardEndpoint(req.params.key);

  console.log(`Forwarding to ${shardEndpoint}`);

  axios.post(shardEndpoint, req.body).then((innerRes) => res.send());

});

app.get('/:key', (req, res) => {

  const shardEndpoint = getShardEndpoint(req.params.key);

  console.log(`Forwarding to ${shardEndpoint}`);

  axios.get(shardEndpoint).then((innerRes) => {

    if (innerRes.data === null) { res.send('null'); return; }

    res.send(innerRes.data);

  });

});

app.listen(8000, () => { console.log('Listening on port 8000');});
```

```
//TO BASICALLY FORWARD DATA TO REVERSE PROXY

const express = require('express');
const fs = require('fs');
const app = express();
const PORT = process.env.PORT;
const DATA_DIR = process.env.DATA_DIR;

app.use(express.json());

app.post('/:key', (req, res) => {
  const { key } = req.params;
  console.log(`Storing data at key ${key}`);
  fs.writeFileSync(destinationFile, req.body.data);
  res.send();
});

app.get('/:key', (req, res) => {
  const { key } = req.params;
  console.log(`Storing data at key${key}`);
  const destinationFile = `${DATA_DIR}/${key}`;
  try { const data = fs.readFileSync(destinationFile); res.send(data); }
  catch (e) { res.send('null'); }
});

app.listen(PORT, () => { console.log(`Listening on port ${PORT}`);});
```

```
replication_and_sharding — node aedb.js — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — node aedb.js
Clements-MBP:replication_and_sharding clementmihai@escu$ DATA_DIR=aedb_data_0 PORT=3000 node aedb.js
Listening on port 3000!
```

```
replication_and_sharding — node aedb.js — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — node aedb.js
Clements-MBP:replication_and_sharding clementmihai@escu$ DATA_DIR=aedb_data_1 PORT=3001 node aedb.js
Listening on port 3001!
Storing data at key a.
```

```
replication_and_sharding — node aedb_proxy.js — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — node aedb_proxy.js
Clements-MBP:replication_and_sharding clementmihai@escu$ node aedb_proxy.js
Listening on port 8000!
Forwarding to: http://localhost:3001/a
```

```
replication_and_sharding — -bash — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — -bash
Clements-MBP:replication_and_sharding clementmihai@escu$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/a
Clements-MBP:replication_and_sharding clementmihai@escu$ curl -w "\n" localhost:8000/a
```

```
replication_and_sharding — node aedb.js — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — node aedb.js
Clements-MBP:replication_and_sharding clementmihai$ DATA_DIR=aedb_data_0 PORT=3000 node aedb.js
Listening on port 3000!
Storing data at key b.
Retrieving data from key b.
Retrieving data from key b.
Storing data at key bar.
Storing data at key baz.
Storing data at key foobar.
Storing data at key foobaz.

replication_and_sharding — node aedb.js — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — node aedb.js
Clements-MBP:replication_and_sharding clementmihai$ DATA_DIR=aedb_data_1 PORT=3001 node aedb.js
Listening on port 3001!
Storing data at key a.
Retrieving data from key a.
Retrieving data from key a.

replication_and_sharding — node aedb_proxy.js — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — node aedb_proxy.js
Clements-MBP:replication_and_sharding clementmihai$ node aedb_proxy.js
Listening on port 8000!
Forwarding to: http://localhost:3001/a
Forwarding to: http://localhost:3001/a
Forwarding to: http://localhost:3000/b
Forwarding to: http://localhost:3000/b
Forwarding to: http://localhost:3001/a
Forwarding to: http://localhost:3000/b
Forwarding to: http://localhost:3000/bar
Forwarding to: http://localhost:3000/baz
Forwarding to: http://localhost:3000/foobar
Forwarding to: http://localhost:3000/foobaz

replication_and_sharding — -bash — 93x24
~/Documents/Content/Design_Fundamentals/Examples/replication_and_sharding — -bash
Clements-MBP:replication_and_sharding clementmihai$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/a
Clements-MBP:replication_and_sharding clementmihai$ curl -w "\n" localhost:8000/a
This is some data.
Clements-MBP:replication_and_sharding clementmihai$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/b
Clements-MBP:replication_and_sharding clementmihai$ curl -w "\n" localhost:8000/b
This is some data.
Clements-MBP:replication_and_sharding clementmihai$ curl -w "\n" localhost:8000/a
This is some data.
Clements-MBP:replication_and_sharding clementmihai$ curl -w "\n" localhost:8000/b
This is some data.
Clements-MBP:replication_and_sharding clementmihai$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/bar
Clements-MBP:replication_and_sharding clementmihai$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/baz
Clements-MBP:replication_and_sharding clementmihai$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/foobar
Clements-MBP:replication_and_sharding clementmihai$ curl --header 'Content-Type: application/json' --data '{"data": "This is some data."}' localhost:8000/foobaz
Clements-MBP:replication_and_sharding clementmihai$ 
```