

PROBLEM STATEMENT

Design a code deployment system for a company that wishes to deploy any web/app/other software services it provides to 100000+ machines spread out across 5-10 regions across the globe in 30 mins tops. (10GB deployed every 30 mins)
Also, you don't need to worry about the testing code aspect and just design for the building and deploying code aspect

CLARIFYING QUES TO ASK

Question 1

Q: What exactly do we mean by a code-deployment system? Are we talking about building, testing, and shipping code?

A: We want to design a system that takes code, builds it into a binary (an opaque blob of data—the compiled code), and deploys the result globally in an efficient and scalable way. We don't need to worry about testing code; let's assume that's already covered.

Question 2

Q: What part of the software-development lifecycle, so to speak, are we designing this for? Is this process of building and deploying code happening when code is being submitted for code review, when code is being merged into a codebase, or when code is being shipped?

A: Once code is merged into the trunk or master branch of a central code repository, engineers should be able to trigger a build and deploy that build (through a UI, which we're not designing). At that point, the code has already been reviewed and is ready to ship. So to clarify, we're not designing the system that handles code being submitted for review or being merged into a master branch—just the system that takes merged code, builds it, and deploys it.

Question 3

Q: Are we essentially trying to ship code to production by sending it to, presumably, all of our application servers around the world?

A: Yes, exactly.

Question 4

Q: How many machines are we deploying to? Are they located all over the world?

A: We want this system to scale massively to hundreds of thousands of machines spread across 5-10 regions throughout the world.

Question 5

Q: This sounds like an internal system. Is there any sense of urgency in deploying this code? Can we afford failures in the deployment process? How fast do we want a single deployment to take?

A: This is an internal system, but we'll want to have decent availability, because many outages are resolved by rolling forward or rolling back buggy code, so this part of the infrastructure may be necessary to avoid certain terrible situations. In terms of failure tolerance, any build should eventually reach a SUCCESS or FAILURE state. Once a binary has been successfully built, it should be shippable to all machines globally within 30 minutes.

Question 6

Q: So it sounds like we want our system to be available, but not necessarily highly available, we want a clear end-state for builds, and we want the entire process of building and deploying code to take roughly 30 minutes. Is that correct?

A: Yes, that's correct.

Question 7

Q: How often will we be building and deploying code, how long does it take to build code, and how big can the binaries that we'll be deploying get?

A: Engineering teams deploy hundreds of services or web applications, thousands of times per day; building code can take up to 15 minutes; and the final binaries can reach sizes of up to 10 GB. The fact that we might be dealing with hundreds of different applications shouldn't matter though; you're just designing the build pipeline and deployment system, which are agnostic to the types of applications that are getting deployed.

Question 8

Q: When building code, how do we have access to the actual code? Is there some sort of reference that we can use to grab code to build?

A: Yes; you can assume that you'll be building code from commits that have been merged into a master branch. These commits have SHA identifiers (effectively arbitrary strings) that you can use to download the code that needs to be built.

ACID transactions will make it safe for potentially hundreds of workers to grab jobs off the queue without unintentionally running the same job twice (we'll avoid race conditions). Our actual transaction will look like this:

```
BEGIN TRANSACTION;
SELECT * FROM jobs_table WHERE status = 'QUEUED' ORDER BY created_at ASC
LIMIT 1;
// if there's none, we ROLLBACK;
UPDATE jobs_table SET status = 'RUNNING' WHERE id = id from previous
query;
COMMIT;
```

The transaction that this auxiliary service will perform will look something like this:

```
UPDATE jobs_table SET status = 'QUEUED' WHERE
status = 'RUNNING' AND
last_heartbeat < NOW() - 10 minutes;
```

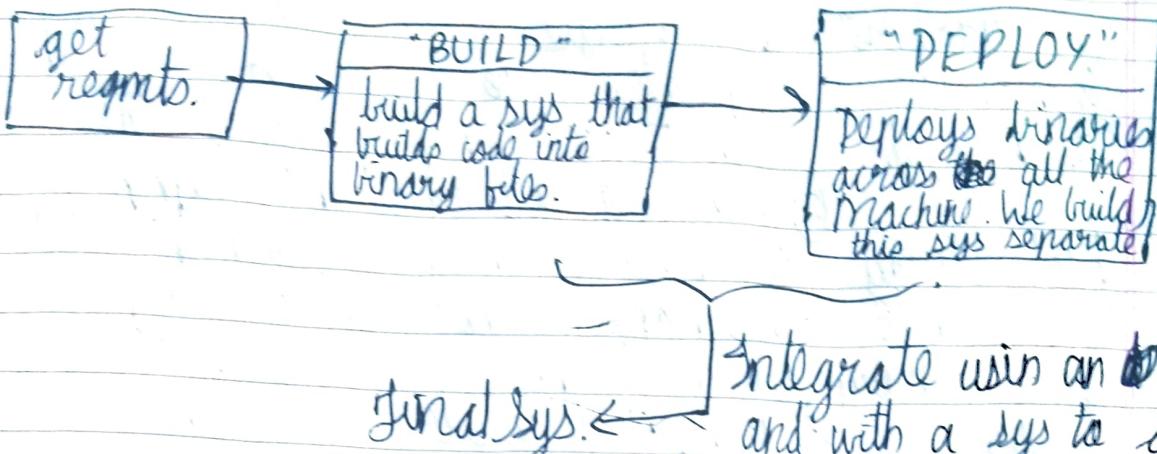
DESIGN-CODE DEPLOYMENT SYSTEM

(A)

General Reqsmts. for sys:

Design	Info	Reason
sys = code deployment sys	Repeatedly (10^3 times/day) builds, deploys code to 10^6 machines across 5-10 regions around the world	
Building code is grabbing snapshots of source code upon SHA identifiers	Implementation details of bldg ac^n is not somethin we need to take care of	eg: we don't need to worry how we build JS or C++ code. Our sys → only enables repeated bldg of code.
Need to split up sys in 2 parts 1) Bldg code 2) Deployin code	→ takes upto 15 mins, ↳ res = binary file of upto 10GB → another 10-15 mins to pass built code to target machines	{ as per accs. we want entire process to happen under 30 mins
Each build will need a clear end state → i) Success or ii) Failure or iii) Running	Availability for sys → 2-3 nines	Don't need to optimize availability as of now

(B) Sys Design Plan



(C) PLAN IMPLEMENT

[I] THE "BUILD" SYS

Decision	Info	Reason
① General Overview	Build sys = queue of "jobs"	Job → process of building code into a binary To ensure proper chronological order in which code is built & eventually deployed
Each job has a commit identifier (SHA identifier)	SHA → secure Hash Algo. converts code into 'encrypted code'	Identifier tells worker what version of code it should build and the name of the "artifact" (name of resulting binary) that'll be created
Assumption → all langs. are auto handled automatically		We don't need to care for type of code being built

Pool of servers known as "workers" to handle all these jobs

Blob storage is chosen to store binary blobs

- ① Each worker takes jobs off the queue repeatedly using FIFO, perhaps
- ② Then, it'll build the binaries (how this happens → not our prob, fn)
- ③ Then it'll write resulting binaries (a.k.a. artifacts) to blob storage (e.g.: Google Cloud storage, Amazon S3)

This way, we ensure that code is built in a chronological order which makes it easier to query it later using its timestamp and thus deployment will also be chronological

Cuz binaries are literally blobs of data so blob storage is art. Binaries need by workers are stored in blob storage in



② JOB QUEUE

SQL dB is chosen

Queue will be implemented as an SQL dB

i) We wanna store historical status in a SQL that, unlike in-mem, persists thru server failures.

Each "job" is a row/record in this SQL table

The attrs/cols per "job" are

- ① id : str : the ID of job, auto generated

- ② created-at : a timestamp as to when job was passed for building

- ③ last-heartbeat

for proper querying of each single job

- ④ commit-SHA : str : the commit SHA identifier

- ⑤ status : str



Date: M T W T F S

idx	blob storage	job's eventual binary in storage			
id	created-at	commit-SHA	name	status	last-heartbeat
job 1					
job 2					
job 3					

general overview of queue implemented as SQL table

NOTE:

- 1) Dequeueing mech. applied by looking at the oldest created-at (timestamp) with a status = QUEUED
- 2) ∵ we gotta idx table at
 - created-at
 - status

Decision	Info	Reason
(3) SQL tables <u>ACID</u> transac ⁿ help ensure no concurrency happens while dequeuing (i.e. no 2 workers tryna grab the same job)	All workers dequeue next job every 5 sec. If we have 100 workers, we are makin $100/5 = 20$ queries to the dB per sec. which is easily handled by an SQL dB	Strong Consistency ensured by ACID transac ⁿ keeps dB updated (i.e. as soon as 1 worker has grabs a job, its status is set to RUNNING so other workers won't grab it)

(4)

Edge Case - Lost Sets

last_heartbeat col. for every job in SQL table. It's basically updates on status of job ~~for~~ every small duration and is stored separately.

Std. time to build code = 15 mins

T/worker updates job's ^{last} heartbeat (of the job it's building) every 3-4 mins i.e.

Updates help us know about the 'health of a server' to deal with cases where server dies mid-build or suffers a network partition or some other failure.

We ~~can~~ use a separate service to continuously monitor the last_heartbeat

If last heartbeat was modified longer than 2 heartbeats before (i.e. 10 mins), something ~~is~~ wrong has happened with server (died, isolated etc.)

↓ this separate service

gets status of the job from RUNNING to QUEUED

(5)

SCALE ESTIMATION

~100 workers reqd. to build code.

No. of builds per day = 5000-10⁴

No. of builds by each worker per day:

$$\frac{1}{\text{day}} = \frac{1}{4 \times 24} = 96$$

$$\frac{\text{No. of builds}}{\text{workers}} = \frac{\text{No. of builds}}{\text{No. of builds by 1 worker}}$$

$$\approx \frac{10000}{96} \approx 100$$

If builds aren't uniformly spread out across the day

i.e. eg:- 8000 builds in day, 2000 at night, we gotta

scale workers horizontally or vertically to handle load.

P2P = peer-to-peer

Date: _____
M W T E S S

⑥. STORAGE

Binaries reqd from workers for each job stored to blob storage

Regional clusters of storage i.e. each region will have its own blob store

Job status col. in SQL table updated to SUCCEEDED

Regional blob stores replicate binaries stored in main blob store asynch.

Worker is released and will now look for the oldest status=QUEUED waala job

faster access to data for each region

II DEPLOYMENT

SYS

Info

Reason

① GENERAL OVERVIEW

3 things reqd.

Peer-to-Peer Network in each region

This P2P network will easily deploy binary stored in this region's blob store very fast

P2P networks help distri. data from 1 storage to several machines real fast

Auxiliary Service 1

Tells us when a binary has been replicated

Once this service tells binary has been replicated!, P2P network starts its work.

Auxiliary Service 2

Serves as a 'source of truth' for what binary should be currly run on all machines

② REPLICAN STATUS SERVICE Auxiliary service

Global service
to check all
regional blob stores
a.k.a. buckets

Once a binary
has been replicated
across all
buckets, this
service updates
a separate SQL
dB where

- ① row = name of
regional bucket
- ② col. 1 = name of
built binary
- ③ col. 2 = replicaⁿ.
status \Rightarrow complete,
failed, pending

Only once a
binary is
replicated across
all buckets i.e.
replicaⁿ status
col. in SQL dB
is completed for
all rows, only
then is it officially
deployable

③ P2P NETWORK

P2P network
relationship
implemented
in regional
clusters i.e.
clusters of
machines in
that region we
 wanna deploy
to

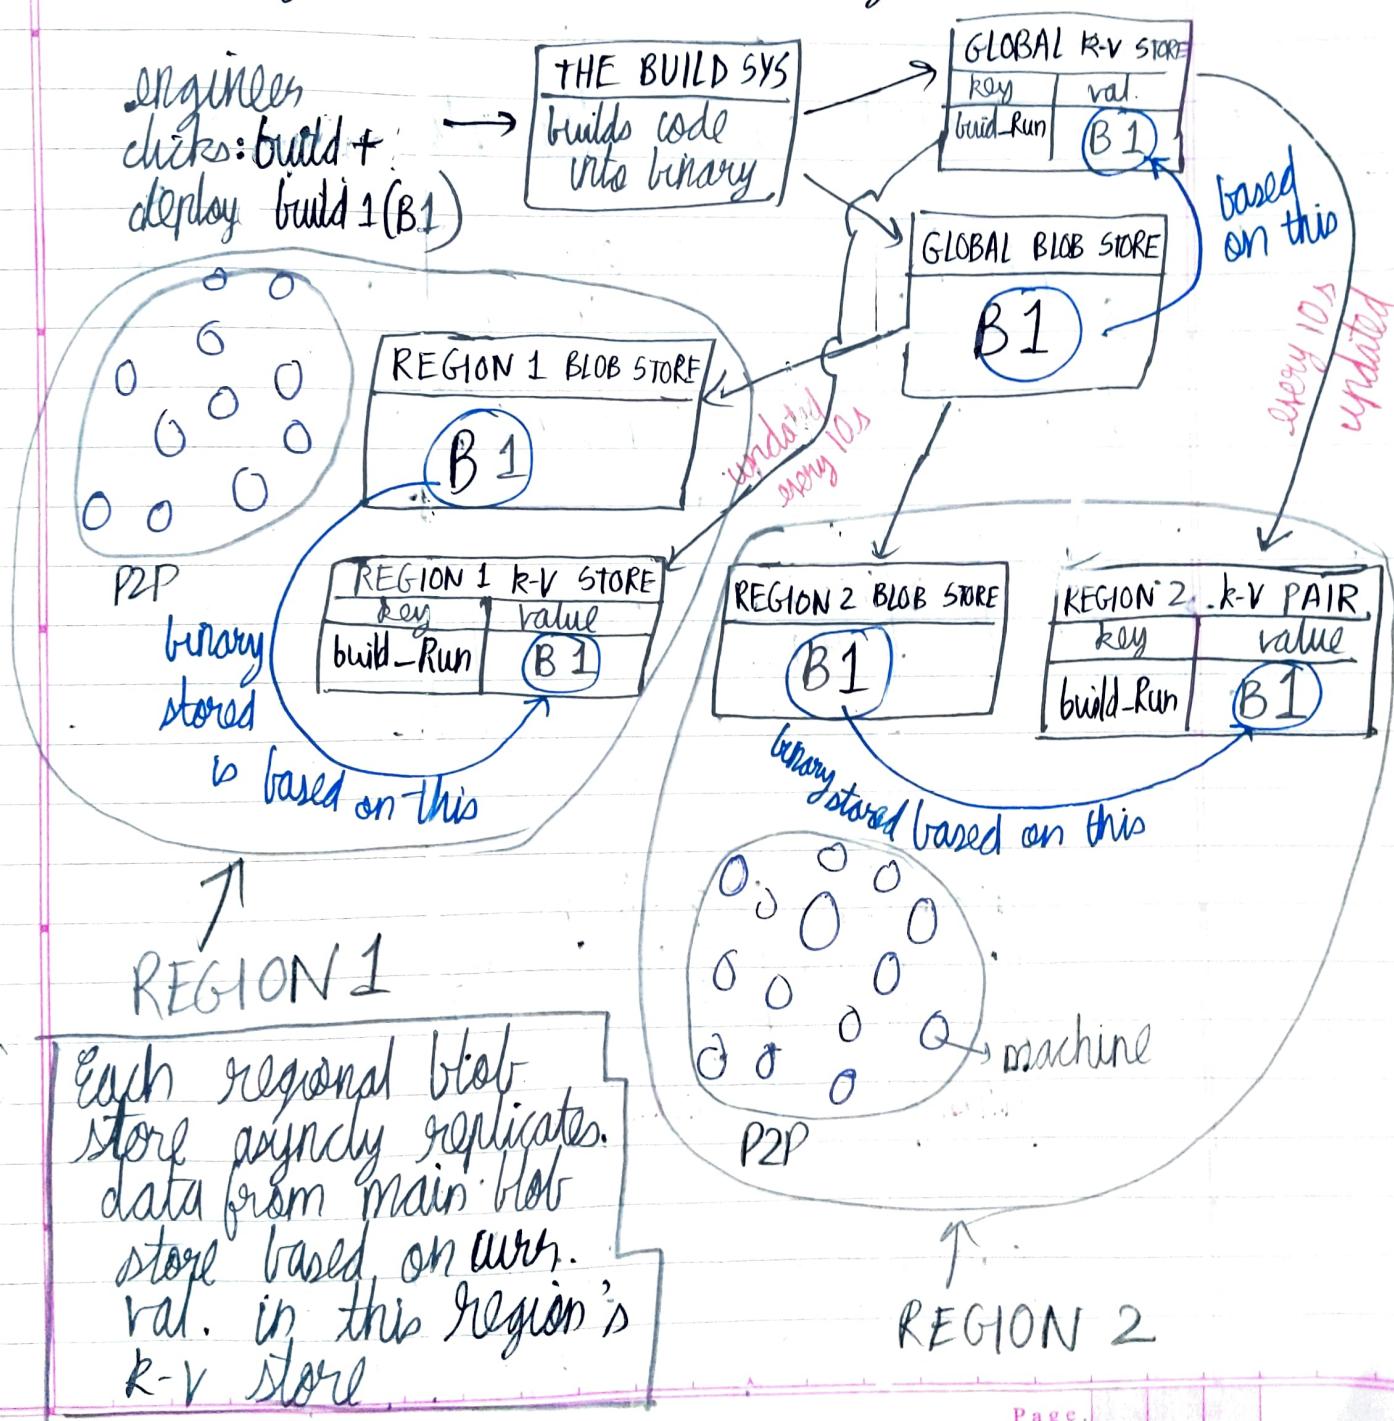
10 GB file is
regional by
↓ regional
P2P
distributes entire
file ~~in~~ in "regional
cluster" very
fast

Auz, P2P networks
are ~~the~~
motherfker
fast.

④ AUXILIARY SERVICE 2 → TRIGGER

Reason → We basically implement this to ensure ~~support~~ support of multiple builds.

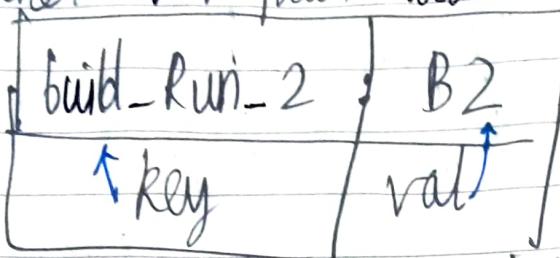
Implementation ⇒ K-V store used for every region as well as our main (a.k.a. "global") K-V store holds config' abt what build should be currently deployed in this region's cluster.



Handling multiple builds

- ① Some other engineer, whilst this build + deploy is ~~not~~ happening, clicks → build + deploy build 2. (B2)

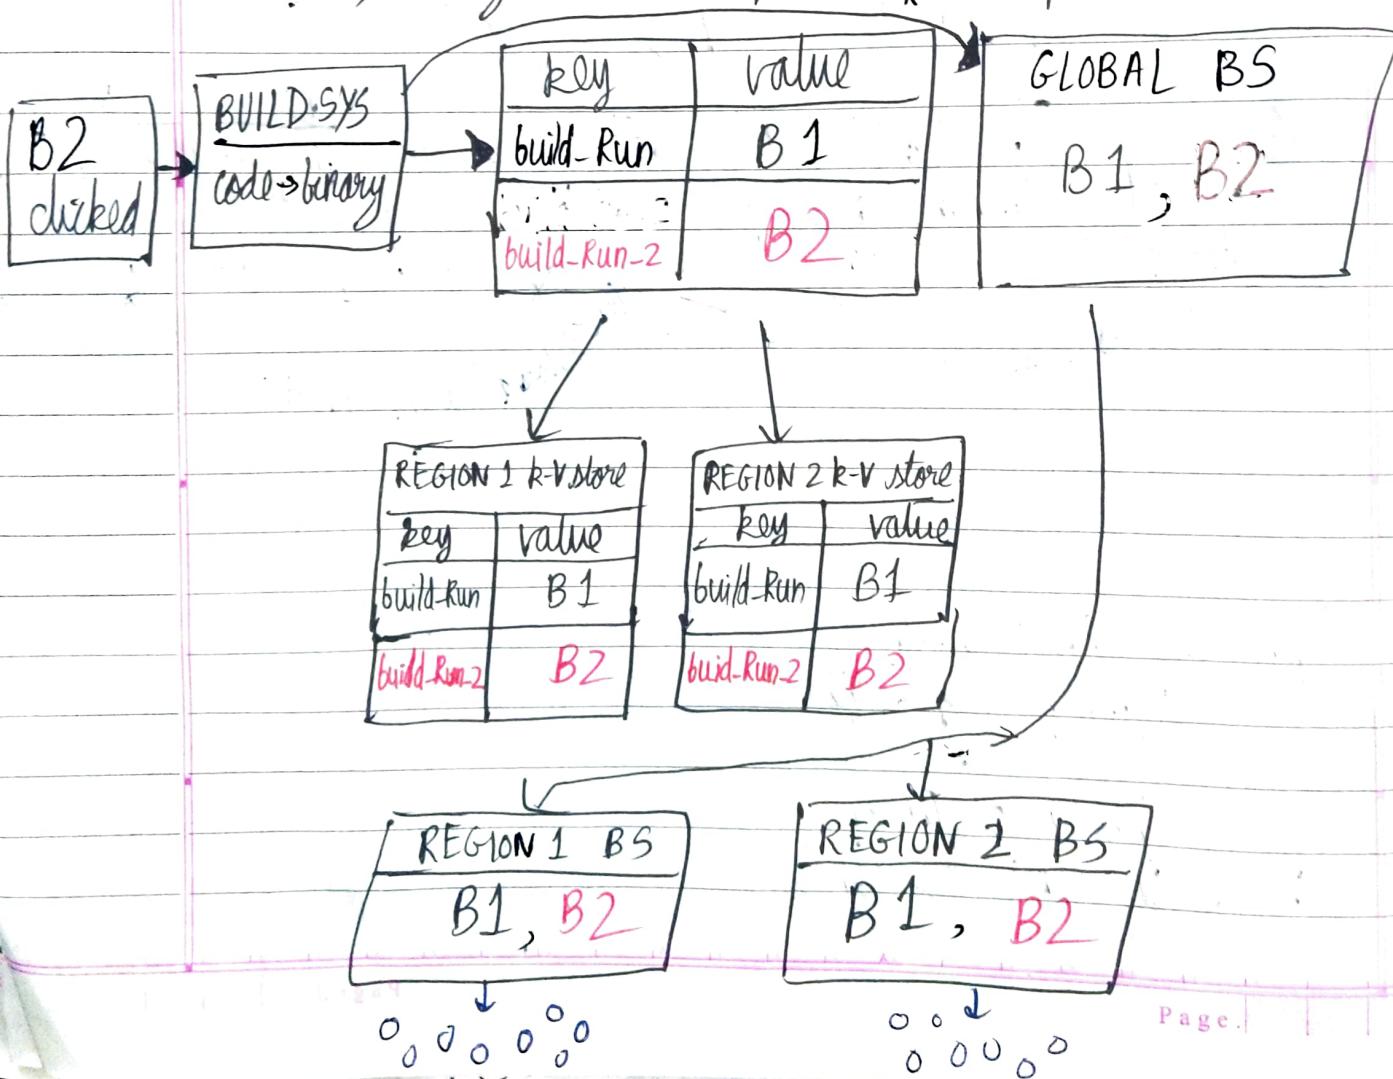
- ② Another k-v pair added in global blob store



- ③ B2 loaded concurrently into global dB

- ④ Every 10s, regional ~~k-v~~ stores query global k-v store for any new k-v pair

Now, they also replicate ^{this} k-v pairs



Now they tell their region's blob store to replicate the binary (B2) as per the KV pair just added

- (5) Each regional BS now replicates B2 binary simult. whilst its say, distributing B1 binary.
- (6) Regional BS deploys new binary B2 to machines via P2P network

